



# OpenGL ES 2.0 Development for the Tegra Platform

---

Version 100106.01

---

# Contents

---

<b>INTRODUCTION</b>	<b>3</b>
<b>SHADER DEVELOPMENT</b>	<b>4</b>
GLSL-ES Shaders	4
Source versus Binary Shaders	4
Precompiling Shaders	5
Loading Shaders	5
Binary Shader-Specific Limitations and Requirements	5
Alpha Blending and Binary Shaders	6
Texture Format and Binary Shaders	9
Framebuffer Format and Binary Shaders	10
Write Masking and Binary Shaders	11
<b>EGL DEVELOPMENT NOTES</b>	<b>11</b>
EGL Configurations	11
<b>TEGRA OPENGL ES LIMITS</b>	<b>13</b>
<b>TEGRA-SUPPORTED OPENGL ES 2.0 EXTENSIONS</b>	<b>14</b>
<b>NVIDIA-PROPRIETARY TEGRA OPENGL ES 2.0 EXTENSION SPECS</b>	<b>16</b>
NV_shader_framebuffer_fetch	16
NV_coverage_sample	19
NV_depth_nonlinear	29
NV_draw_path	32
NV_system_time	47

## Introduction

This documentation provides platform-specific details for developing shader-based OpenGL ES 2.0 (GLES2) applications on the Tegra platform. The information in this documentation is designed to be OS-independent, and represents the capabilities of the Tegra OpenGL ES 2.0 hardware and driver on all supported operating systems.

This document does not detail performance optimizations; it only covers rendering feature- and compatibility-related items. Performance optimizations for Tegra are described in other documentation that may be available from the Tegra Developers' site.

## Shader Development

This section provides guidelines and details of shader development on the Tegra platform. It discusses specific shader features and limitations on the Tegra as well as differences between shaders supplied to GLES2 as source code and as precompiled binaries. It also provides information on the shader compiler tools required to generate precompiled binary shaders for the Tegra.

### GLSL-ES Shaders

The Tegra supports OpenGL ES 2.0 and its shading language, GLSL-ES. Basically a subset of desktop GLSL, GLSL-ES removes all of the fixed-function language constructs, and also removes language constructs for GL features that are not a part of OpenGL ES 2.0 core, such as 1D and 3D textures.

General GLSL-ES features and uses are outside the scope of this document. Developers should refer directly to the GLSL specification, which is currently downloadable from the Khronos group website:

[http://www.khronos.org/registry/gles/specs/2.0/GLSL\\_ES\\_Specification\\_1.0.14.pdf](http://www.khronos.org/registry/gles/specs/2.0/GLSL_ES_Specification_1.0.14.pdf)

Obviously, developers can also refer to applications and shaders provided online and in NVIDIA's Tegra Applications SDK for basic examples of GLSL-ES.

### Source versus Binary Shaders

As a mobile platform, the Tegra GLES2 drivers are focused on optimal performance and memory footprint while maximizing the feature set. The driver supports both source code shaders provided to GLES2 at runtime, as well as platform-specific binary shaders compiled offline specifically for Tegra. Applications may choose one or the other path (or even a mixture of both) as desired. Precompiled shaders provide the highest performance at runtime, as they can be used directly by the hardware. However, there are limitations and hints that must be provided when using precompiled shaders that will be detailed in other sections of this document. Source code shaders, on the other hand, are extremely flexible and can even be generated as strings on the fly in application code. However, the driver must compile (and in some cases recompile) these shaders at runtime on the Tegra itself, increasing memory footprint and incurring performance overhead at points when the shaders must be recompiled. This can include recompiling the shader during rendering if a never-before-seen combination of some key rendering states is used. The states that can cause shader recompilation are the same as those listed later in this document as requiring pragmas when used in binary, precompiled shaders.

Both shaders provided to GLES2 as source at runtime and precompiled shaders can use the same shader source code in almost all cases. But there are significant cases in which the two may need to differ, which will be detailed in other sections of this documentation.

## Precompiling Shaders

The generation of binary, precompiled shaders is dependent upon the host PC OS (generally Windows or Linux). The resulting shaders can be OS-image-specific, although these incompatibilities are avoided whenever possible. See the documentation supplied with the particular Platform Support pack for details on compiling shaders and the resulting binary shaders.

## Loading Shaders

### Source Code Shaders

Shaders are loaded using the OpenGL ES standard functions: `glShaderSource`, `glCompileShader`, and `glLinkProgram`.

### Binary Shaders

Shaders are loaded as binaries on the Tegra platform via the OpenGL ES standard functions: `glShaderBinary` and `glLinkProgram`. The shader format enumerant is listed in `gl2ext.h`, and is `GL_NVIDIA_PLATFORM_BINARY_NV`.

## Binary Shader-Specific Limitations and Requirements

This section includes information on binary, precompiled shader limitations and requirements. In order to optimally support binary shaders on Tegra, the compiled shader binary needs to be specialized for several bits of OpenGL-ES state. In particular, these states are:

**Alpha Blending** : OpenGL-ES fixed-function alpha blending is performed inside the shader processor on Tegra. To support this, the compiler must accept the blend state and perform the specified blend equation inside the compiled fragment program. Simply setting the alpha blending state in the application C code will not suffice. A pragma in the shader code or a command-line argument to the shader compiler is required.

**Texture format** : Tegra supports unsigned fixed-point, signed fixed-point, thin floating-point and wide-floating point texture formats. Depending upon which texture format is used for a given sampler, the compiler may need to perform different post-fetch operations, or assume different return value registers. Thus, if you are using floating-point or signed textures ("classic" OpenGL texture types are generally unsigned fixed-point), a pragma in the shader code or a command-line argument to the shader compiler is required. Tegra supports up to 16 simultaneous samplers.

**Framebuffer format(s)** : When reading from or writing to the framebuffer, the compiler will need to be aware of whether the framebuffer format is fixed- or floating- point. If the framebuffer format is floating-point, a pragma in the shader code or a command-line argument to the shader compiler is required. Tegra supports up to 8 simultaneous color targets (multiple render targets, a.k.a. MRT).

**Write Masking** : OpenGL-ES allows applications to selectively enable writes to individual color components. To support this, the shader compiler must be given the write mask; the setting it in the C-code GL calls will not suffice. A pragma in the shader code or a command-line argument to the shader compiler is required.

These required pragmas and tags ***DO NOT apply to shaders provided to GLES2 as source code***, as the driver can recompile the shader as needed for the specific fixed-function settings.

## Alpha Blending and Binary Shaders

When used with binary shaders, Tegra requires that alpha blending modes be built into the binary shaders. The alpha blending mode as set by the C code calls to OpenGL ES will not set the blending mode on Tegra.

Binary shaders on Tegra do not use the OpenGL ES `glBlend` mode settings to determine blending. The blending mode must be compiled into the shader.

The Tegra blending mode can be set in one of two ways. One method involves an addition to the shader source itself (a pragma), while the other involves an additional compiler flag be passed to the shader compiler.

In the case of the shader source addition, the programmer simply adds a pragma to the top of the fragment shader source code that declares the alpha blending modes to be used. This is convenient to specify, but does have the drawback of “burning” the alpha blending mode into the shader source itself. Precompiled, that shader will always use the given blending mode. But the pragma is ignored in shaders supplied to GLES2 at runtime as source code, with one exception: Any calls to `glValidateShader` will test the pragma against the current blend settings.

The compiler flag option allows for the same source shader code to be compiled into multiple binary shaders with different blending modes. However, it requires a more advanced content pipeline on the part of the app, which can determine the required settings when compiling the shaders.

In order to ensure that source code shaders and precompiled binary shaders match, the blending mode burned into the shader should match the OpenGL ES shading mode set in C code at the time of each draw call. If these modes do not match, calls to `glValidateShader` will fail. Also, the draw call itself may generate a GL error (but the geometry will be drawn based on the shader’s burned-in blending mode).

### Alpha Blending Pragma/Compile Option Syntax

The syntax of both the pragma and the compile time option are equivalent, and basically follow the OpenGL ES blending mode enumerations.

Pragma Syntax:

```
#pragma profile
#pragma blendoperation( <drawbufferName>, <rgbEquation>,
                      <rgbSrcFactor>, <rgbDstFactor>, [ <alphaEquation>,
                      <alphaSrcFactor>, <alphaDstFactor>] )
```

### Compile Time Syntax:

```
-profileopt -
blendoperation, <drawbufferName>, <rgbEquation>, <rgbSrcFactor>,
<rgbDstFactor>[, <alphaEquation>, <alphaSrcFactor>, <alphaDstFactor>]
```

<drawbufferName> should be the name of a draw buffer that is written in the shader, i.e.:  
`gl_FragColor` or `gl_FragData[0]` – `gl_FragData[7]`. The name of the semantics used to describe each output is also accepted, i.e. `COLOR` or `COLOR0-COLOR7` (`COL` and `COL0-COL7` are also accepted).

<rgbEquation> and <alphaEquation> may be one of:

```
GL_NO_OPERATION
GL_FUNC_ADD
GL_FUNC_SUBTRACT
GL_FUNC_REVERSE_SUBTRACT
```

If <rgbEquation> or <alphaEquation> is `GL_NO_OPERATION`, no operation is performed.

If <rgbEquation> or <alphaEquation> is `GL_FUNC_ADD`, alpha blending performs the operation:

```
<rgbEquation>  dst.rgb = <rgbSrcFactor>*src.rgb + <rgbDstFactor>*dst.rgb
<alphaEquation> dst.a  = <alphaSrcFactor>*src.a + <alphaDstFactor>*dst.a
```

If <rgbEquation> or <alphaEquation> is `GL_FUNC_SUBTRACT`, alpha blending performs the operation:

```
<rgbEquation>  dst.rgb = <rgbSrcFactor>*src.rgb - <rgbDstFactor>*dst.rgb
<alphaEquation> dst.a  = <alphaSrcFactor>*src.a - <alphaDstFactor>*dst.a
```

Finally, if <rgbEquation> or <alphaEquation> is `GL_FUNC_REVERSE_SUBTRACT`, alpha blending performs the operation:

```
<rgbEquation>  dst.rgb = <rgbSrcFactor>*dst.rgb - <rgbDstFactor>*src.rgb
```

$\text{dst.a} = \text{alphaSrcFactor} * \text{dst.a} - \text{alphaDstFactor} * \text{src.a}$

`<rgbDstFactor>` and `<alphaDstFactor>` may be one of:

GL\_ZERO  
GL\_ONE  
GL\_SRC\_COLOR  
GL\_ONE\_MINUS\_SRC\_COLOR  
GL\_DST\_COLOR  
GL\_ONE\_MINUS\_DST\_COLOR  
GL\_SRC\_ALPHA  
GL\_ONE\_MINUS\_SRC\_ALPHA  
GL\_DST\_ALPHA  
GL\_ONE\_MINUS\_DST\_ALPHA  
GL\_CONSTANT\_COLOR  
GL\_ONE\_MINUS\_CONSTANT\_COLOR  
GL\_CONSTANT\_ALPHA  
GL\_ONE\_MINUS\_CONSTANT\_ALPHA

### Blending Pragma Examples

Some examples of pragmas and the resulting blend modes include:

“Standard” alpha-based blending:

```
#pragma profile
#pragma blendoperation(gl_FragColor, GL_FUNC_ADD,
    GL_SRC_ALPHA,
    GL_ONE_MINUS_SRC_ALPHA)
```

Color modulation:

```
#pragma profile
#pragma blendoperation(gl_FragColor, GL_FUNC_ADD,
    GL_ZERO,
    GL_SRC_COLOR)
```

Color add:

```
#pragma profile
#pragma blendoperation(gl_FragColor, GL_FUNC_ADD,
    GL_ONE, GL_ONE)
```



Alpha blend the colors, add the alphas:

```
#pragma profile
#pragma blendoperation( gl_FragColor, GL_FUNC_ADD,
                       GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA,
                       GL_FUNC_ADD, GL_ONE, GL_ONE)
```

## Texture Format and Binary Shaders

Pragma Syntax:

```
#pragma profile
#pragma samplerformat( <samplerName>: <format> )
```

Configuration File Syntax:

```
-profileopts -samplerformat,<samplerName>,<format>
```

<samplerName> should be the name of a sampler uniform appearing inside the shader. <format> should be one of:

```
fixed_unsigned (default)
fixed_signed
half_float_luminance_alpha
half_float_packed
half_float
```

fixed\_unsigned corresponds to hardware texture formats of:

```
A8, I8, L8, L8A8, B2G3R3, B5G6R5, B5G5R5A1, B4G4R4A4, A1B5G5R5,
A4B4G4R4,
R8G8B8A8, B8G8R8A8, DXT1, DXT1C, DXT3, DXT5, ETC, LATC1 and LATC2
```

fixed\_signed corresponds to software texture formats of:

```
COMPRESSED_SIGNED_LUMINANCE_ALPHA_LATC1_EXT (hardware: LATC1) and
COMPRESSED_SIGNED_LUMINANCE_ALPHA_LATC2_EXT (hardware: LATC2)
```

half\_float\_luminance\_alpha corresponds to hardware texture formats of:

```
A16_float, L16_float and L16A16_float
```

half\_float\_packed corresponds to hardware texture formats of:

```
R11G11B10_float
```

half\_float corresponds to hardware texture formats of:

```
R16G16B16A16_float
```

If no format is specified for a sampler, a default format of `fixed_unsigned` is assumed. If a format is specified for a sampler not appearing in the shader, a warning is generated. If <format> is not one of

the supported formats, a warning is generated, and compilation proceeds as if the format was specified as `fixed_unsigned`.

### Texture Format Pragma Examples

```
#pragma profilepragma samplerformat( bumpmap: fixed_signed )
#pragma profilepragma samplerformat( specularCube: half_float_packed )
// all other samplers default to fixed_unsigned
```

## Framebuffer Format and Binary Shaders

Pragma Syntax:

```
#pragma profilepragma drawbufferformat( <drawbufferName>: <format> )
```

Configuration File Syntax:

```
-profileopts -drawbufferformat,<drawbufferName>,<format>
```

`<drawbufferName>` should be the name of a draw buffer that is written in the shader, i.e.: `gl_FragColor` or `gl_FragData[0]` – `gl_FragData[7]`. The name of the semantics used to describe each output is also accepted, i.e. `COLOR` or `COLOR0` – `COLOR7` (`COL` and `COL0`–`COL7` are also accepted). `<format>` should be one of:

```
fixed_unsigned (default)
half_float_luminance_alpha
half_float_packed
half_float
```

The list of hardware formats mirrors the list of texture formats. `fixed_signed` is not supported, since Tegra does not support any renderable signed formats.

### Framebuffer Pragma Examples

```
#pragma profilepragma drawbufferformat(gl_FragData[3]:
half_float_luminance_alpha)
#pragma profilepragma drawbufferformat(gl_FragData[0]:
half_float_packed)
// any other gl_FragData outputs default to fixed_unsigned
```

## Write Masking and Binary Shaders

Currently, on Tegra with binary shaders, the write mask set by `glColorMask` is not sufficient to enable write masking; the write mask set via this function must be matched with pragmas in the shader.

Binary shaders on Tegra using writemasking must match the OpenGL ES `glColorMask` mode with a masking mode pragma that must be compiled into the shader. `glValidateProgram` can be used to verify this matching state.

Pragma Syntax:

```
#pragma profilepragma colorwritemask(<drawbuffername>, <red>, <green>,  
<blue>, <alpha> )
```

Configuration File Syntax:

```
profileopts -  
colorwritemask,<drawbuffername>,<red>,<green>,<blue>,<alpha>
```

<drawbufferName> should be the name of a draw buffer that is written in the shader, i.e.: `gl_FragColor` or `gl_FragData[0] - gl_FragData[7]`. The name of the semantics used to describe each output is also accepted, i.e. `COLOR` or `COLOR0- COLOR7` (`COL` and `COL0-COL7` are also accepted).

Legal values for <red>, <green>, <blue> and <alpha> are:

false or true

Write-masking behaves identically to the definition in section 4.2.2 of the OpenGL 2.0 Specification (page 214). A value of "true" indicates that the resulting color value (after blending, if necessary) will be stored in the framebuffer. A value of "false" indicates that the original framebuffer color value should be maintained.

If unspecified, the default value for `colorwritemask` is `TRUE, TRUE, TRUE, TRUE`.

## EGL Development Notes

### EGL Configurations

EGL's method of sorting configurations returned from queries is often counter-intuitive. Applications using `eglChooseConfig` **must not** simply select the first returned configuration, nor should they request only one configuration. An example of a common and confusing case is requesting a 565 RGB configuration. Owing to section 3.4 of the EGL spec, `eglChooseConfig` must return the **deepest**

color buffer first, even if it is deeper than the requested format, and ***even if the requested format could have been matched exactly***. In other words, an implementation that supports 565 and 8888 must return 8888 earlier in the list than 565, even if 565 is requested. The EGL spec notes the following in a footnote to 3.4:

“This rule places configs with deeper color buffers first in the list returned by `eglChooseConfig`. Applications may find this counterintuitive, and need to perform additional processing on the list of configs to find one best matching their requirements. For example, specifying RGBA depths of 5651 could return a list whose first config has a depth of 8888.”

The side-effects of this can be detrimental in subtle ways. For example, if an application requests 565 but uses the first returned configuration blindly, on Tegra they will receive an 8888 configuration. Rendering to an 8888 back buffer:

- 1) Can require additional memory bandwidth above rendering to 565
- 2) Will silently disable dithering (as there is no need for it)

If this 8888 backbuffer is used with a 565 onscreen surface (the default on most OS images), then the 8888 surface will be converted to 565 during the backbuffer-to-frontbuffer copy. This format-converting copy is not dithered, and can result in far worse color banding than would be seen with a correctly dithered 565 backbuffer direct-copied into a 565 frontbuffer. Simply searching the returned configurations for a 565 configuration can increase performance and visual quality!

Applications should always request an array of multiple configurations, and should query important attributes such as red, green and blue depths of each, performing their own manual sorting and filtering of the resulting array. EGL's behavior is defined by the spec; Tegra's driver cannot deviate from the proscribed order.

## Tegra OpenGL ES Limits

The following table lists the current limits of various OpenGL ES 2.0 values as returned by the driver.

Note that these are intended as guidelines – applications should always query important values from the particular driver being used.

GL_SUBPIXEL_BITS	4
GL_ALIASED_POINT_SIZE_RANGE	(1, 256)
GL_ALIASED_LINE_WIDTH_RANGE	(1, 256)
GL_MAX_ARRAY_TEXTURE_LAYERS_EXT	2048
GL_MAX_CUBE_MAP_TEXTURE_SIZE	2048
GL_MAX_TEXTURE_SIZE	2048
GL_MAX_TEXTURE_MAX_ANISOTROPY_EXT	15
GL_MAX_VIEWPORT_DIMS	(3839, 3839)
GL_NUM_COMPRESSED_TEXTURE_FORMATS	9
GL_NUM_SHADER_BINARY_FORMATS	1
GL_SHADER_BINARY_FORMATS	GL_NVIDIA_PLATFORM_BINARY_NV
GL_MAX_VERTEX_ATTRIBS	16
GL_MAX_VERTEX_UNIFORM_VECTORS	256
GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS	0
GL_MAX_VARYING_VECTORS	15
GL_MAX_TEXTURE_IMAGE_UNITS	16
GL_MAX_FRAGMENT_UNIFORM_VECTORS	1024
GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS	16
GL_MAX_COLOR_ATTACHMENTS_NV	8
GL_MAX_RENDERBUFFER_SIZE	3839
GL_MAX_DRAW_BUFFERS_ARB	8

## Tegra-Supported OpenGL ES 2.0 Extensions

The following are extensions supported by the OpenGL ES 2.0 drivers in most Tegra OS images.

However, this list is not universal. Applications should call `glGetString(GL_EXTENSIONS)` to verify the support for key extensions. The first group of extensions are GL- or OES-standard, and the documentation can be found at <http://www.opengl.org/> or <http://www.khronos.org/>. Notes are made only as needed.

- `GL_ARB_draw_buffers`
- `GL_ARB_half_float_pixel`
- `GL_EXT_packed_float`
- `GL_EXT_texture_array`
- `GL_EXT_texture_compression_latc`
- `GL_EXT_texture_filter_anisotropic`
- `GL_OES_compressed_ETC1_RGB8_texture`
- `GL_OES_EGL_image`
- `GL_OES_fbo_render_mipmap`
- `GL_OES_shader_binary`
  - (Indicated by the value of `GL_NUM_SHADER_BINARY_FORMATS` being nonzero). Indicates that the implementation supports precompiled binary shaders. All of the demos use this capability on Tegra. See the `nv_shader` helper library for details
- `GL_OES_texture_float`
- `GL_OES_vertex_half_float`
- `GL_EXT_texture_compression_dxt1`
  - The implementation supports specifying textures with the `GL_COMPRESSED_RGBA_S3TC_DXT1_EXT` formats. Not exported on Tegra, but supported
- `GL_EXT_texture_compression_s3tc`
  - The implementation supports specifying textures with the `GL_COMPRESSED_RGBA_S3TC_DXT[1,3,5]_EXT` formats.
- `GL_OES_framebuffer_object`
  - (Required extension.) Framebuffer objects are supported. Not exported as an extension string on Tegra, but supported
- `GL_OES_mapbuffer`
  - The implementation supports the `glMapBufferOES` and `glUnmapBufferOES` functions. These are exported directly and do not need to be queried.
- `GL_OES_rgb8_rgba8`
  - The implementation supports `GL_RGBA8_OES` and `GL_RGB8_OES` as FBO color buffer formats.
- `GL_OES_stencil8`

- Indicates that the implementation can support an 8-bit stencil buffer for render targets. Not exported as an extension string on Tegra, but supported.
- GL\_OES\_texture\_half\_float
  - Indicates that the GL\_HALF\_FLOAT\_OES token is accepted by `glTex[Sub]Image2D` and allows for the creation and use of 16-bit floating-point textures (1 sign bit, 5 exponent bits, 10 mantissa bits). These are supported as texture formats and FBO formats. An example of loading a half-float RGBA texture is:

```
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGBA, width, height, 0, GL_RGBA, GL_HALF_FLOAT_OES, img);
```

The following extensions are NVIDIA-specific; their extension specs (as available) are at the end of this chapter.

- GL\_NV\_coverage\_sample
- GL\_NV\_depth\_nonlinear
- GL\_NV\_fbo\_color\_attachments
- GL\_NV\_read\_buffer
- GL\_NV\_draw\_path

# NVIDIA-Proprietary Tegra OpenGL ES 2.0 Extension Specs

## NV\_shader\_framebuffer\_fetch

### Name

NV\_shader\_framebuffer\_fetch

### Name Strings

GL\_NV\_shader\_framebuffer\_fetch

### Contact

Gary King, NVIDIA Corporation (gking 'at' nvidia.com)

### Notice

Copyright NVIDIA Corporation, 2005 - 2006

### Status

NVIDIA Proprietary

### Version

Last Modified: 2006/04/28  
NVIDIA Revision: 0.9

### Number

XXXX Not Yet XXXX

### Dependencies

This extension is written against the OpenGL-ES Shading Language 1.10 Specification.

### Overview

This extension provides a mechanism whereby a fragment shader written in the OpenGL-ES Shading Language (GLSL-ES) may use the color values stored in the active draw buffers as well-defined input values.

### Issues

1. How should this functionality be exposed?

RESOLVED: Four options were considered for this functionality:

- A) Defining reads from gl\_FragColor and gl\_FragData (prior to any writes) to result in the existing framebuffer values.
- B) Defining new read-only built-in variables corresponding to the existing framebuffer data (e.g., gl\_LastFragColor).



- C) Defining new built-in functions which return the existing framebuffer data.
- D) Defining a new programmable stage (e.g., a Sample Shader) which takes the fragment shader output values and the existing framebuffer data as inputs.

This extension has chosen option B, as it provides the best mix of language / API simplicity and programmer flexibility. Reusing the existing built-in variables (option (A)) unnecessarily complicates the language, since it requires shader compilers to perform flow analysis to determine whether or not framebuffer loads are required. Option (C) requires adding either multiple entry points (one for each `gl_FragData` array entry), or also adding keywords to specify which buffer should be read. Option (D) is interesting, but leads to a variety of questions regarding what functionality should be included in the new stage (e.g., "should texturing be included?") Therefore, Option B has been chosen due to its comparative simplicity and available functionality.

2. What value should `gl_LastFragColor` contain when the `ARB_draw_buffers` extension is in use?

RESOLVED: Two options were considered for this functionality:

- A) To mirror `ARB_draw_buffers` specification the value in `gl_LastFragColor` should be dependent on all active draw buffers.
- B) The value on `gl_LastFragColor` should be the same as the one in `gl_LastFragData[0]`.

This extension has chosen option B, for its simplicity. Option A may follow the spirit of the original `ARB_draw_buffers` specification more closely but the value of `gl_LastFragData` becomes somewhat undefined. There is no good way to combine the values from multiple draw buffers.

#### New Functions and Entry Points

None

#### New Builtin Variables

```
mediump vec4 gl_LastFragData[gl_MaxDrawBuffers]
mediump vec4 gl_LastFragColor
```

#### Changes to the OpenGL-ES Shading Language 1.10 Specification, Chapter 7

Add after the last sentence of Paragraph 2 of Section 7.2, Fragment Shader Special Variables ("These variables may be written to more [...]":

"... To access the existing framebuffer values (e.g., to implement a complex blend operation inside the shader), fragment shaders should use the read-only input values `gl_LastFragColor` and `gl_LastFragData`."

Insert new paragraph after Paragraph 7 of Section 7.2 ("If a shader statically assigns [...]")

"Similarly, if a shader using the NV\_shader\_framebuffer\_fetch extension statically assigns a value to gl\_FragColor, it may not read any element of gl\_LastFragData. If a shader using the NV\_shader\_framebuffer\_fetch extension statically writes a value to any element of gl\_FragData, it may not read from gl\_LastFragColor. That is, use of the inputs defined in the NV\_shader\_framebuffer\_fetch extension must mirror the outputs used in the shader program."

## NV\_coverage\_sample

### Name

NV\_coverage\_sample

### Name Strings

GL\_NV\_coverage\_sample  
EGL\_NV\_coverage\_sample

### Contact

Gary King, NVIDIA Corporation (gking 'at' nvidia.com)

### Notice

Copyright NVIDIA Corporation, 2005 - 2007

### Status

NVIDIA Proprietary

### Version

Last Modified Date: 2007/03/20  
NVIDIA Revision: 1.0

### Number

XXXX Not Yet XXXX

### Dependencies

Written based on the wording of the OpenGL 2.0 specification  
and the EXT\_framebuffer\_object specification.

Written based on the wording of the EGL 1.2 specification.

Requires OpenGL-ES 2.0 and OES\_framebuffer\_object.

Requires EGL 1.1.

### Overview

Anti-aliasing is a critical component for delivering high-quality OpenGL rendering. Traditionally, OpenGL implementations have implemented two anti-aliasing algorithms: edge anti-aliasing and multisampling.

Edge anti-aliasing computes fractional fragment coverage for all primitives in a rendered frame, and blends edges of abutting and/or overlapping primitives to produce smooth results. The image quality produced by this approach is exceptionally high; however, applications are render their geometry perfectly ordered back-to-front in order to avoid artifacts such as bleed-through. Given the algorithmic complexity and performance cost of performing exact geometric sorts, edge anti-aliasing has been used very sparingly, and almost never in interactive games.

Multisampling, on the other hand, computes and stores subpixel (a.k.a. "sample") coverage for rasterized fragments, and replicates all post-alpha test operations (e.g., depth test, stencil test, alpha blend) for each sample. After the entire scene is rendered, the samples are filtered to compute the final anti-aliased image. Because the post-alpha test operations are replicated for each sample, all of the bleed-through and ordering artifacts that could occur with edge anti-aliasing are avoided completely; however, since each sample must be computed and stored separately, anti-aliasing quality is limited by framebuffer storage and rendering performance.

This extension introduces a new anti-aliasing algorithm to OpenGL, which dramatically improves multisampling quality without adversely affecting multisampling's robustness or significantly increasing the storage required, coverage sampling.

Coverage sampling adds an additional high-precision geometric coverage buffer to the framebuffer, which is used to produce high-quality filtered results (with or without the presence of a multisample buffer). This coverage information is computed and stored during rasterization; since applications may render objects where the specified geometry does not correspond to the visual result (examples include alpha-testing for "imposters," or extruded volume rendering for stencil shadow volumes), coverage buffer updates may be masked by the application, analogous to masking the depth buffer.

#### IP Status

NVIDIA Proprietary

#### New Procedures and Functions

```
void CoverageMaskNV( boolean mask )
void CoverageOperationNV( enum operation )
```

#### New Tokens

Accepted by the <attrib\_list> parameter of eglChooseConfig and eglCreatePbufferSurface, and by the <attribute> parameter of eglGetConfigAttrib

(Note: these enumerants reuse the values GLX\_VIDEO\_OUT\_DEPTH\_NV - GLX\_VIDEO\_OUT\_COLOR\_AND\_ALPHA\_NV from the GLX\_NV\_video\_out extension)

```
EGL_COVERAGE_BUFFERS_NV          0x20C5
EGL_COVERAGE_SAMPLES_NV         0x20C6
```

(Note: these enumerants reuse the values REGISTER\_COMBINERS\_NV - VARIABLE\_G\_NV of the NV\_register\_combiners extension)

Accepted by the <internalformat> parameter of RenderbufferStorageEXT and the <format> parameter of ReadPixels

```
COVERAGE_COMPONENT_NV           0x8522
```

Accepted by the <internalformat> parameter of RenderbufferStorageEXT

```
COVERAGE_COMPONENT4_NV          0x8523
```

Accepted by the <operation> parameter of CoverageOperationNV

COVERAGE_ALL_FRAGMENTS_NV	0x8524
COVERAGE_EDGE_FRAGMENTS_NV	0x8525
COVERAGE_AUTOMATIC_NV	0x8526

Accepted by the <attachment> parameter of  
FramebufferRenderbuffer, and GetFramebufferAttachmentParameteriv

COVERAGE_ATTACHMENT_NV	0x8527
------------------------	--------

Accepted by the <buf> parameter of Clear

COVERAGE_BUFFER_BIT_NV	0x8000
------------------------	--------

Accepted by the <pname> parameter of GetIntegerv

COVERAGE_BUFFERS_NV	0x8528
COVERAGE_SAMPLES_NV	0x8529

#### Changes to Chapter 4 of the OpenGL 2.0 Specification

Insert a new section, after Section 3.2.1 (Multisampling)

##### "3.2.2 Coverage Sampling

Coverage sampling is a mechanism to antialias all GL primitives: points, lines, polygons, bitmaps and images. The technique is similar to multisampling, with all primitives being sampled multiple times at each pixel, and a sample resolve applied to compute the color values stored in the framebuffer's color buffers. As with multisampling, coverage sampling resolves color sample and coverage values to a single, displayable color each time a pixel is updated, so antialiasing appears to be automatic at the application level. Coverage sampling may be used simultaneously with multisampling; however, this is not required.

An additional buffer, called the coverage buffer, is added to the framebuffer. This buffer stores additional coverage information that may be used to produce higher-quality antialiasing than what is provided by conventional multisampling.

When the framebuffer includes a multisample buffer (3.5.6), the samples contain this coverage information, and the framebuffer does not include the coverage buffer.

If the value of COVERAGE\_BUFFERS\_NV is one, the rasterization of all primitives is changed, and is referred to as coverage sample rasterization. Otherwise, primitive rasterization is referred to as multisample rasterization (if SAMPLE\_BUFFERS is one) or single-sample rasterization (otherwise). The value of COVERAGE\_BUFFERS\_NV is queried by calling GetIntegerv with <pname> set to COVERAGE\_BUFFERS\_NV.

During coverage sample rasterization the pixel fragment contents are modified to include COVERAGE\_SAMPLES\_NV coverage values. The value of COVERAGE\_SAMPLES\_NV is an implementation-dependent constant, and is queried by calling GetIntegerv with <pname> set to COVERAGE\_SAMPLES\_NV.

The command

```
CoverageOperationNV(enum operation)
```

may be used to modify the manner in which coverage sampling is performed for all primitives. If <operation> is COVERAGE\_ALL\_FRAGMENTS\_NV, coverage sampling will be performed and the coverage buffer updated for all fragments generated during rasterization. If <operation> is COVERAGE\_EDGE\_FRAGMENTS\_NV, coverage sampling will only be performed for fragments generated at the edge of the primitive (by only updating fragments at the edges of primitives, applications may get better visual results when rendering partially transparent objects). If <operation> is COVERAGE\_AUTOMATIC\_NV, the GL will automatically select the appropriate coverage operation, dependent on the GL blend mode and the use of gl\_LastFragColor / gl\_LastFragData in the bound fragment program. If blending is enabled, or gl\_LastFragColor / gl\_LastFragData appears in the bound fragment program, COVERAGE\_AUTOMATIC\_NV will behave identically to COVERAGE\_EDGE\_FRAGMENTS\_NV; otherwise, COVERAGE\_AUTOMATIC\_NV will behave identically to COVERAGE\_ALL\_FRAGMENTS\_NV. The default coverage operation is COVERAGE\_AUTOMATIC\_NV."

Insert a new section, after Section 3.3.3 (Point Multisample Rasterization)

#### "3.3.4 Point Coverage Sample Rasterization

If the value of COVERAGE\_BUFFERS\_NV is one, then points are rasterized using the following algorithm, regardless of whether point antialiasing (POINT\_SMOOTH) is enabled or disabled. Point rasterization produces fragments using the same algorithm described in section 3.3.3; however, sample points are divided into SAMPLES multisample points and COVERAGE\_SAMPLES\_NV coverage sample points.

Rasterization for multisample points uses the algorithm described in section 3.3.3. Rasterization for coverage sample points uses implementation-dependent algorithms, ultimately storing the results in the coverage buffer."

Insert a new section, after Section 3.4.4 (Line Multisample Rasterization)

#### "3.4.5 Line Coverage Sample Rasterization

If the value of COVERAGE\_BUFFERS\_NV is one, then lines are rasterized using the following algorithm, regardless of whether line antialiasing (LINE\_SMOOTH) is enabled or disabled. Line rasterization produces fragments using the same algorithm described in section 3.4.4; however, sample points are divided into SAMPLES multisample points and COVERAGE\_SAMPLES\_NV coverage sample points.

Rasterization for multisample points uses the algorithm described in section 3.4.4. Rasterization for coverage sample points uses implementation-dependent algorithms, ultimately storing results in the coverage buffer."

Insert a new section, after Section 3.5.6 (Polygon Multisample Rasterization)

#### "3.5.7 Polygon Coverage Sample Rasterization

If the value of COVERAGE\_BUFFERS\_NV is one, then polygons are rasterized using the following algorithm, regardless of whether polygon antialiasing (POLYGON\_SMOOTH) is enabled or disabled. Polygon rasterization produces fragments using the same algorithm described in section 3.5.6; however, sample points are divided into SAMPLES multisample points and COVERAGE\_SAMPLES\_NV coverage sample points.

Rasterization for multisample points uses the algorithm described in section 3.5.7. Rasterization for coverage sample points uses implementation-dependent algorithms, ultimately storing results in the coverage buffer."

Insert a new section, after Section 3.6.6 (Pixel Rectangle Multisample Rasterization)

### "3.6.7 Pixel Rectangle Coverage Sample Rasterization

If the value of COVERAGE\_BUFFERS\_NV is one, then pixel rectangles are rasterized using the algorithm described in section 3.6.6."

Modify the first sentence of the second-to-last paragraph of section 3.7 (Bitmaps) to read:

### "Bitmap Multisample and Coverage Sample Rasterization

If MULTISAMPLE is enabled, and the value of SAMPLE\_BUFFERS is one; or if the value of COVERAGE\_BUFFERS\_NV is one, then bitmaps are rasterized using the following algorithm. [...]"

Insert after the first paragraph of Section 4.2.2 (Fine Control of Buffer Updates):

"The coverage buffer can be enabled or disabled for writing coverage sample values using

```
void CoverageMaskNV( boolean mask );
```

If <mask> is non-zero, the coverage buffer is enabled for writing; otherwise, it is disabled. In the initial state, the coverage buffer is enabled for writing."

And change the text of the last 2 paragraphs of Section 4.2.2 to read:

"The state required for the various masking operations is three integers and two bits: an integer for color indices, an integer for the front and back stencil values, a bit for depth values, and a bit for coverage sample values. A set of four bits is also required indicating which components of an RGBA value should be written. In the initial state, the integer masks are all ones, as are the bits controlling the depth value, coverage sample value and RGBA component writing.

### Fine Control of Multisample Buffer Updates

When the value of SAMPLE\_BUFFERS is one, ColorMask, DepthMask, CoverageMask, and StencilMask or StencilMaskSeparate control the modification of values in the multisample buffer. [...]"

Change paragraph 2 of Section 4.2.3 (Clearing the Buffers) to read:

"is the bitwise OR of a number of values indicating which buffers are to be cleared. The values are COLOR\_BUFFER\_BIT, DEPTH\_BUFFER\_BIT, STENCIL\_BUFFER\_BIT, ACCUM\_BUFFER\_BIT and COVERAGE\_BUFFER\_BIT\_NV, indicating the buffers currently enabled for color writing, the depth buffer, the stencil buffer, the accumulation buffer and the virtual-coverage buffer, respectively. [...]"

Insert a new paragraph after paragraph 4 of Section 4.3.2 (Reading Pixels) (beginning with "If there is a multisample buffer ..."):

"If the <format> is COVERAGE\_COMPONENT\_NV, then values are taken from the coverage buffer; again, if there is no coverage buffer, the error INVALID\_OPERATION occurs. When <format> is COVERAGE\_COMPONENT\_NV, <type> must be GL\_UNSIGNED\_BYTE. Any other value for <type> will generate the error INVALID\_ENUM. If there is a multisample buffer, the values are undefined."

#### Modifications to the OES\_framebuffer\_object specification

Add a new table at the end of Section 4.4.2.1 (Renderbuffer Objects)

Sized internal format	Base Internal Format	C Samples
COVERAGE_COMPONENT4_NV	COVERAGE_COMPONENT_NV	4

Table 1.000 Desired component resolution for each sized internal format that can be used only with renderbuffers"

Add to the bullet list in Section 4.4.4 (Framebuffer Completeness)

"An internal format is 'coverage-renderable' if it is COVERAGE\_COMPONENT\_NV or one of the COVERAGE\_COMPONENT\_NV formats from table 1.000. No other formats are coverage-renderable"

Add to the bullet list in Section 4.4.4.1 (Framebuffer Attachment Completeness)

"If <attachment> is COVERAGE\_ATTACHMENT\_NV, then <image> must have a coverage-renderable internal format."

Add a paragraph at the end of Section 4.4.4.2 (Framebuffer Completeness)

"The values of COVERAGE\_BUFFERS\_NV and COVERAGE\_SAMPLES\_NV are derived from the attachments of the currently bound framebuffer object. If the current FRAMEBUFFER\_BINDING\_OES is not 'framebuffer-complete', then both COVERAGE\_BUFFERS\_NV and COVERAGE\_SAMPLES\_NV are undefined. Otherwise, COVERAGE\_SAMPLES\_NV is equal to the number of coverage samples for the image attached to COVERAGE\_ATTACHMENT\_NV, or zero if COVERAGE\_ATTACHMENT\_NV is zero."

#### Additions to the EGL 1.2 Specification

Add to Table 3.1 (EGLConfig attributes)

Attribute	Type	Notes



EGL_COVERAGE_BUFFERS_NV	integer	number of coverage buffers	
EGL_COVERAGE_SAMPLES_NV	integer	number of coverage samples per	
		pixel	
+-----+-----+-----+-----+			

Modify the first sentence of the last paragraph of the "Buffer Descriptions and Attributes" subsection of Section 3.4 (Configuration Management), p. 16

"There are no single-sample depth, stencil or coverage buffers for a multisample EGLConfig; the only depth, stencil and coverage buffers are those in the multisample buffer. [...]"

And add the following text at the end of that paragraph:

"The <coverage buffer> is used only by OpenGL ES. It contains primitive coverage information that is used to produce a high-quality anti-aliased image. The format of the coverage buffer is not specified, and its contents are not directly accessible. Only the existence of the coverage buffer, and the number of coverage samples it contains, are exposed by EGL.

EGL\_COVERAGE\_BUFFERS\_NV indicates the number of coverage buffers, which must be zero or one. EGL\_COVERAGE\_SAMPLES\_NV gives the number of coverage samples per pixel; if EGL\_COVERAGE\_BUFFERS\_NV is zero, then EGL\_COVERAGE\_SAMPLES\_NV will also be zero."

Add to Table 3.4 (Default values and match criteria for EGLConfig attributes)

+-----+-----+-----+-----+-----+					
Attribute	Default	Selection	Sort	Sort	
		Criteria	Order	Priority	
+-----+-----+-----+-----+-----+					
EGL_COVERAGE_BUFFERS_NV	0	At Least	Smaller	7	
EGL_COVERAGE_SAMPLES_NV	0	At Least	Smaller	8	
+-----+-----+-----+-----+-----+					

And renumber existing sort priorities 7-11 as 9-13.

Modify the list in "Sorting of EGLConfigs" (Section 3.4.1, pg 20)

```
" [...]
5. Smaller EGL_SAMPLE_BUFFERS
6. Smaller EGL_SAMPLES
7. Smaller EGL_COVERAGE_BUFFERS_NV
8. Smaller EGL_COVERAGE_SAMPLES_NV
9. Smaller EGL_DEPTH_SIZE
10. Smaller EGL_STENCIL_SIZE
11. Smaller EGL_ALPHA_MASK_SIZE
12. Special: [...]
13. Smaller EGL_CONFIG_ID [...]"
```

#### Usage Examples

##### (1) Basic Coverage Sample Rasterization

```
glCoverageMaskNV(GL_TRUE);
glDepthMask(GL_TRUE);
glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE);

while (1)
```

```

{
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT |
           GL_COVERAGE_BUFFER_BIT_NV);
    glDrawElements(...);
    eglSwapBuffers(...);
}

```

## (2) Multi-Pass Rendering Algorithms

```

while (1)
{
    glDepthMask(GL_TRUE);
    glCoverageMaskNV(GL_TRUE);
    glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE);
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT |
           GL_COVERAGE_BUFFER_BIT_NV);

    // first render pass: render Z-only (occlusion surface), with
    // coverage info. color writes are disabled

    glCoverageMaskNV(GL_TRUE);
    glDepthMask(GL_TRUE);
    glColorMask(GL_FALSE, GL_FALSE, GL_FALSE, GL_FALSE);
    glDepthFunc(GL_LESS);
    glDrawElements(...);

    // second render pass: set Z test to Z-equals, disable Z-writes &
    // coverage writes. enable color writes. coverage may be
    // disabled, because subsequent rendering passes are rendering
    // identical geometry -- since the final coverage buffer will be
    // unchanged, we can disable coverage writes as an optimization.

    glCoverageMaskNV(GL_FALSE);
    glDepthMask(GL_FALSE);
    glDepthFunc(GL_EQUAL);
    glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE);
    glDrawElements(...);

    eglSwapBuffers();
}

```

## (3) Rendering Translucent Objects on Top of Opaque Objects

```

while (1)
{
    glDepthMask(GL_TRUE);
    glCoverageMaskNV(GL_TRUE);
    glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE);
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT |
           GL_COVERAGE_BUFFER_BIT_NV);

    // render opaque, Z-buffered geometry with coverage info for the
    // entire primitive. Overwrite coverage data for all fragments, so
    // that interior fragments do not get resolved incorrectly.

    glDepthFunc(GL_LESS);
    glCoverageOperationNV(GL_COVERAGE_ALL_FRAGMENTS_NV);
    glDrawElements(...);

    // render translucent, Z-buffered geometry. to ensure that visible

```

```

// edges of opaque geometry remain anti-aliased, change the
// coverage operation to just edge fragments. this will maintain
// the coverage information underneath the translucent geometry,
// except at translucent edges.

glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);
glCoverageOperationNV(GL_COVERAGE_EDGE_FRAGMENTS_NV);
glEnable(GL_BLEND);
glDrawElements(...);
glDisable(GL_BLEND);

eglSwapBuffers();
}

```

(4) Rendering Opacity-Mapped Particle Systems & HUDs on Top of Opaque Geometry

```

while (1)
{
    glDepthMask(GL_TRUE);
    glCoverageMaskNV(GL_TRUE);
    glColorMask(GL_TRUE, GL_TRUE, GL_TRUE, GL_TRUE);
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT |
           GL_COVERAGE_BUFFER_BIT_NV);

    // render opaque, Z-buffered geometry, with coverage info.
    glDepthFunc(GL_LESS);
    glDrawElements(...);

    // render opacity-mapped geometry. disable Z writes, enable alpha
    // blending. also, disable coverage writes -- the edges of the
    // geometry used for the HUD/particle system have alpha values
    // tapering to zero, so edge coverage is uninteresting, and
    // interior coverage should still refer to the underlying opaque
    // geometry, so that opaque edges visible through the translucent
    // regions remain anti-aliased.

    glCoverageMaskNV(GL_FALSE);
    glDepthMask(GL_FALSE);
    glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);
    glEnable(GL_BLEND);
    glDrawElements(...);
    glDisable(GL_BLEND);

    eglSwapBuffers();
}

```

#### Issues

1. Is any specific discussion of coverage sampling resolves required, particularly with respect to application-provided framebuffer objects?

RESOLVED: No. Because the coverage sampling resolve is an implementation-dependent algorithm, it is always legal behavior for framebuffer read / copy functions to return the value in the selected ReadBuffer as if COVERAGE\_BUFFERS\_NV was zero. This allows textures attached to the color attachment points of framebuffer objects to behave predictably, even when COVERAGE\_BUFFERS\_NV is one.

Implementations are encouraged, whenever possible, to use the highest-quality coverage sample resolve supported for calls to `eglSwapBuffers`, `eglCopyBuffers`, `ReadPixels`, `CopyPixels` and `CopyTex{Sub}Image`.

2. Should all render buffer & texture types be legal sources for image resolves and coverage attachment?

RESOLVED: This spec should not place any arbitrary limits on usage; however, there are many reasons why implementers may not wish to support coverage sampling for all surface types.

Implementations may return `FRAMEBUFFER_UNSUPPORTED_OES` from `CheckFramebufferStatusOES` if an object bound to `COVERAGE_ATTACHMENT_NV` is incompatible with one or more objects bound to `DEPTH_ATTACHMENT_OES`, `STENCIL_ATTACHMENT_OES`, or `COLOR_ATTACHMENTi_OES`.

#### Revision History

#1.0 - 20.03.2007

Renumbered enumerants. Reformatted to 80 columns.

## NV\_depth\_nonlinear

### Name

NV\_depth\_nonlinear

### Name Strings

GL\_NV\_depth\_nonlinear  
EGL\_NV\_depth\_nonlinear

### Contact

Gary King, NVIDIA Corporation (gking 'at' nvidia.com)

### Notice

Copyright NVIDIA Corporation, 2005 - 2007.

### Status

NVIDIA Proprietary

### Version

Last Modified: 2007/03/20  
NVIDIA Revision: 1.0

### Number

XXXX Not Yet XXXX

### Dependencies

Written based on the wording of the OpenGL 2.0 Specification and  
EGL 1.2 Specification.

Requires EGL 1.1.

Requires OpenGL-ES 1.0.

OES\_framebuffer\_object affects the wording of this specification.

### Overview

Due to the perspective divide, conventional integer Z-buffers have a hyperbolic distribution of encodings between the near plane and the far plane. This can result in inaccurate depth testing, particularly when the number of depth buffer bits is small and objects are rendered near the far plane.

Particularly when the number of depth buffer bits is limited (desirable and/or required in low-memory environments), artifacts due to this loss of precision may occur even with relatively modest far plane-to-near plane ratios (e.g., greater than 100:1).

Many attempts have been made to provide alternate encodings for Z-buffer (or alternate formulations for the stored depth) to reduce the artifacts caused by perspective division, such as W-buffers, Z-complement buffers and floating-point 1-Z buffers.

This extension adds a non-linear encoded Z buffer to OpenGL, which can improve the practically useful range of, e.g. 16-bit depth buffers by up to a factor of 16, greatly improving depth test quality in applications where the ratio between the near and far planes can not be as tightly controlled.

#### IP Status

NVIDIA Proprietary

#### New Procedures and Functions

None

#### New Tokens

(Note: these enumerant values reuse the GLX\_VIDEO\_OUT\_COLOR\_NV - GLX\_VIDEO\_OUT\_ALPHA\_NV values from the GLX\_NV\_video\_out extension, and the COMBINER3\_NV value from the NV\_register\_combiners extension)

Accepted as a valid sized internal format by all functions accepting sized internal formats with a base format of DEPTH\_COMPONENT

DEPTH\_COMPONENT16\_NONLINEAR\_NV      0x8553

Accepted by the <attrib\_list> parameter of eglChooseConfig and eglCreatePbufferSurface, and by the <attribute> parameter of eglGetConfigAttrib

EGL\_DEPTH\_ENCODING\_NV      0x20C3

Accepted as a value in the <attrib\_list> parameter of eglChooseConfig and eglCreatePbufferSurface, and returned in the <value> parameter of eglGetConfigAttrib

EGL\_DEPTH\_ENCODING\_NONE\_NV      0  
EGL\_DEPTH\_ENCODING\_NONLINEAR\_NV      0x20C4

#### Changes to the OpenGL 2.0 Specification

Add the following line to table 3.16 (p. 154)

Sized Internal Format	Base Internal Format	D
		Bits
DEPTH_COMPONENT16_NONLINEAR_NV	DEPTH_COMPONENT	16

#### Changes to the EGL 1.2 Specification

Add the following line to table 3.1 (p. 14)

Attribute	Type	Notes
EGL_DEPTH_ENCODING_NV	enum	Type of depth-buffer encoding employed

Modify the description of the depth buffer in Section 3.4 (p. 15)

"The depth buffer is used only by OpenGL ES. It contains fragment depth (Z) information generated during rasterization. EGL\_DEPTH\_SIZE indicates the depth of this buffer in bits, and EGL\_DEPTH\_ENCODING\_NV indicates which alternate depth-buffer encoding (if any) should be used. Legal values for EGL\_DEPTH\_ENCODING\_NV are: EGL\_DONT\_CARE, EGL\_DEPTH\_ENCODING\_NONE\_NV and EGL\_DEPTH\_ENCODING\_NONLINEAR\_NV."

Add the following line to table 3.4 (p. 20)

Attribute	Default	Selection	Sort	Sort
		Criteria	Order	Priority
EGL_DEPTH_ENCODING_NV	EGL_DONT_CARE	Exact	None	-

#### Issues

None

#### Revision History

#1.0 - 20.03.2007

Renumbered enumerants. Reformatted to 80 columns.

## NV\_draw\_path

### Name

NV\_draw\_path

### Name Strings

GL\_NV\_draw\_path

### Contact

Jussi Rasanen, NVIDIA Corporation (jrasanen 'at' nvidia.com)  
Tero Karras, NVIDIA Corporation (tkarras 'at' nvidia.com)

### Notice

Copyright NVIDIA Corporation, 2008

### Status

NVIDIA Proprietary

### Version

Last Modified: 2008/09/16  
NVIDIA Revision: 0.11

### Number

XXXX Not Yet XXXX

### Dependencies

Written based on the wording of the OpenGL 2.0 Specification.

Requires OpenGL-ES 2.0.

### Overview

This extension adds functionality to render planar Bezier paths. Use cases for this extension include acceleration of vector graphics content and text rendering.

A path is defined as a number of segments, representing either straight lines, or quadratic or cubic Bezier curves, and can be either filled or stroked. Filling corresponds to generating the fragments that lie within the interior of the path. Stroking corresponds to generating the fragments that lie within a region defined by sweeping a straight-line pen along the path.

Path segments are specified using a command array and a vertex coordinate array. When a path is drawn, the command array is processed sequentially. There are two categories of commands: ones that cause a path segment to be drawn, and ones that affect how path is rendered. Depending on its type, each command consumes a variable number of coordinates from the vertex coordinate array.

When filling a path, the order in which the path segments are specified is disregarded. The only requirement is that they form zero or more closed



contours. If a path contains unclosed contours, its interior and thus the resulting set of fragments is undefined.

The contours of a path may have self-intersecting geometry and overlap with each other. For such paths, the interior is determined using a fill rule. Two fill rules, even-odd and non-zero, are provided. The direction of path segments matters only with the non-zero fill rule, as explained below.

When stroking a path, additional cap and join styles may be applied at the start and end of path segments. Joins are automatically generated between pairs of segments whose corresponding commands are adjacent. Caps are generated based on explicit path commands.

Rendering quality can be controlled per path by specifying the maximum deviation from the ideal curve in window space.

#### Path rendering pipeline

The path rendering pipeline consists of three stages: transformation and texture coordinate generation, fill and stroke rasterization, and fragment shader. This extension provides a minimal fixed function transformation and texture coordinate generation stage. Programmable vertex shaders are not supported in the context of path rendering.

#### Path definition

Paths are defined as a combination of an immutable sequence of commands and an associated mutable sequence of vertex coordinates. Each command consumes zero or more vertex coordinates. Path commands are represented as unsigned bytes, whereas the data type of the vertex coordinates is specified separately for each path. Path vertices are always two-dimensional.

The following table lists the available path commands:

Path command	Coords	Notes
MOVE_TO_NV	2	Change the current position
LINE_TO_NV	2	Draw a straight line
QUADRATIC_BEZIER_TO_NV	4	Draw a quadratic Bezier curve
CUBIC_BEZIER_TO_NV	6	Draw a cubic Bezier curve
START_MARKER_NV	0	Record the current position
CLOSE_NV	0	Draw line to the recorded position
STROKE_CAP0_NV	0	Use cap style 0 in adjacent segment
STROKE_CAP1_NV	0	Use cap style 1 in adjacent segment
STROKE_CAP2_NV	0	Use cap style 2 in adjacent segment
STROKE_CAP3_NV	0	Use cap style 3 in adjacent segment

#### Transformation and texture coordinate generation

Path vertices specified by the vertex coordinate sequence are converted to the homogenous form (x, y, 0, 1) by the transformation stage, and then transformed from model space to clip space using the MATRIX\_PATH\_TO\_CLIP\_NV matrix.

To facilitate texture mapping and color gradients, the path vertices are also transformed using each of the MATRIX\_PATH\_COORD[0-3]\_NV matrices. A

built-in fragment shader varying array `gl_PathCoord` of type `vec4` receives the corresponding interpolated values. The number of elements in the `gl_PathCoord` array is 4. Although gradients and texture coordinates can also be implemented using the `gl_FragCoord` built-in fragment shader variable, it is generally more efficient to use `gl_PathCoord`, avoiding unnecessary per-fragment matrix multiplications.

`MATRIX_PATH_COORD[0-3]_NV` and `MATRIX_PATH_TO_CLIP_NV` can define a homogenous perspective transformation. It is up to the fragment shader to normalize the interpolated coordinates if necessary.

#### Filling a path

When filling a path, the path segments must form zero or more closed contours. If any of the contours are left open, the resulting set of fragments is undefined. This requirement can be rephrased as follows, depending on the value of `FILL_RULE_NV`:

##### `NON_ZERO_NV`:

Each two-dimensional point has an equal number of path segments starting and ending at it.

##### `EVEN_ODD_NV`

Each two-dimensional point has an even number of path segments starting or ending at it.

Note that for any two points to be considered identical, the binary representations of their coordinates must match exactly.

To determine the segments to draw, the path commands are processed sequentially. The following temporary values are maintained during the process:

`i`: Current vertex coordinate index, initially 0.

`cp`: Current position, initially (0, 0).

`sp`: Start position, initially undefined.

Each path command is processed depending on its type as follows. `c[i]` is used to denote the *i*'th value in the vertex coordinate array.

##### `MOVE_TO_NV`:

Replace the current position.  
`cp = (c[i+0], c[i+1]), i += 2.`

##### `LINE_TO_NV`:

Draw a straight line from `<cp>` to `(c[i+0], c[i+1])`.  
`cp = (c[i+0], c[i+1]), i += 2.`

##### `QUADRATIC_BEZIER_TO_NV`:

Draw a quadratic Bezier curve from `<cp>` to `(c[i+2], c[i+3])` using `(c[i+0], c[i+1])` as the control point.  
`cp = (c[i+2], c[i+3]), i += 4.`

##### `CUBIC_BEZIER_TO_NV`:

Draw a cubic Bezier curve from `<cp>` to `(c[i+4], c[i+5])` using `(c[i+0], c[i+1])` and `(c[i+2], c[i+3])` as the control points.  
`cp = (c[i+4], c[i+5]), i += 6.`

##### `START_MARKER_NV`:

Replace the start position.

sp = cp.

**CLOSE\_NV:**

If <sp> is undefined, ignore the command.  
Otherwise, draw a straight line from <cp> to <sp>.  
cp = sp.

**STROKE\_CAP[0-3]\_NV:**

Ignore the command.

START\_MARKER\_NV and CLOSE\_NV commands can be used to implement subpath closure found in many vector graphics content formats. For filled paths, an explicit LINE\_TO\_NV command to the start position will produce the same result as CLOSE\_NV. For stroked paths, the difference is that CLOSE\_NV will join the closing line segment to the segment following the START\_MARKER\_NV command.

A fill rule is applied to determine if any given point is contained within the interior of the path. The fill rules are defined by projecting a ray from the point in question to infinity and counting the intersections of the ray and path segments. When looking along the direction of the ray, segments intersecting from left to right increment the counter and right to left segment intersections decrement the counter. If the fill rule is NON\_ZERO\_NV, the point is within the interior if the final count is non-zero. If the fill rule is EVEN\_ODD\_NV, the point is within the interior if the final count is odd. The counter must support at least 255 intersections. For more complex paths, the results are undefined.

Curves may be approximated within a limit specified by the PATH\_QUALITY\_NV parameter. The limit defines the radius of a disc in the window space. Placing the disc at each sampling point, the following rules are used to determine whether to generate the corresponding fragments:

- \* If the disc is entirely inside the path, generate a fragment.
- \* If the disc is entirely outside the path, do not generate a fragment.
- \* If the disc is partially inside the path, whether to generate a fragment is up to the implementation.

**Stroking a path**

Stroking is performed by sweeping a straight-line pen along each path segment, generating fragments for the sampling points touched by the pen. Additionally, cap and join styles may be applied at the start and end of the segments.

Cap and join styles are selected for each path segment based on the path commands adjacent to the one specifying the segment, and the values of the path parameters. The general rule is that the end of a segment is joined to start of the following segment if they are specified by adjacent path commands. If the start or end of a segment is not joined, a cap is generated instead.

Fill rule is not applied when stroking. Instead, a fragment is generated for each sampling point inside the stroke. Even in case the stroke sweeps over a sampling point multiple times, only one fragment is generated.

Dashing is not supported directly. Instead, this extension allows implementing dashing in user code by generating the corresponding.

Paths are stroked in a coordinate space distinct from the path user space

and the clip space. The transformation from the stroke space to the path user space is controlled by the `MATRIX_STROKE_TO_PATH_NV` matrix. Both the path user space and the stroke space are two-dimensional, and thus only the 2x2 upper-left components of the matrix are used.

Conceptually, stroking a path consists of five steps. First, the path segments are transformed from the path user space to the stroke space using the inverse of the stroke-to-path matrix. Second, the set of points affected by the stroke is determined in the stroke space, using a straight-line pen that extends one unit into each direction. Third, the set of points is transformed from the stroke space back to the path user space using the stroke-to-path matrix. Fourth, the points are further transformed from the path user space to the clip space using the path-to-clip matrix. Fifth, a fragment is generated for each sampling point contained by the set of transformed points.

The stroke-to-path matrix allows specifying stroke width independent of how the path itself is transformed. Two common scenarios include scaling stroke, where the stroke width varies as the path-to-clip transformation changes, and non-scaling stroke, where the width remains constant in the clip space. For scaling stroke, the stroke-to-path matrix should be specified as an identity matrix multiplied by half of the desired stroke width in the path user space. For non-scaling stroke, it should be specified as the inverse of the path-to-clip matrix multiplied by half of the stroke width in the clip space.

The style of all joins is determined by the `STROKE_JOIN_STYLE_NV` path parameter, which can be set to one of the following values:

`JOIN_MITER_NV:`

Extend the incoming and outgoing stroke outlines until they intersect. If the distance between the intersection point and the center point exceeds `STROKE_MITER_LIMIT_NV` in the stroke space, apply a bevel join instead.

`JOIN_ROUND_NV:`

Connect the incoming and outgoing stroke outlines with a circular arc segment in the stroke space, corresponding to a radius of one unit.

`JOIN_BEVEL_NV:`

Connect the incoming and outgoing stroke outlines with a straight line.

`JOIN_CLIPPED_MITER_NV:`

Same as `JOIN_MITER_NV` if `STROKE_MITER_LIMIT_NV` is not exceeded. Otherwise, clip the extended outlines and connect them with a straight line. The clipping is done against a line whose distance from the center point is equal to `STROKE_MITER_LIMIT_NV` in the stroke space, and whose orientation is symmetrical with regards to the outlines.

The style of a start cap depends on the previous path command, and the style of an end cap depends on the next command. If the command is not `STROKE_CAP[0-3]_NV`, the cap style is `STROKE_CAP_BUTT_NV`. Otherwise, the style is determined by the corresponding `STROKE_CAP[0-3]_STYLE_NV` path parameter, each of which can be set to one of the following values:

`CAP_BUTT_NV:`

Terminate the segment with a straight line connecting the two outline endpoints.

**CAP\_ROUND\_NV:**

Terminate the segment with a semicircle with radius equal to one in the stroke space.

**CAP\_SQUARE\_NV:**

Terminate the segment with a rectangle extending one unit along the path tangent.

**CAP\_TRIANGLE\_NV:**

Terminate the segment with a triangle with two vertices at the stroke outline endpoints, and a third vertex one unit along the path tangent.

As with fill, the path commands are processed sequentially, maintaining the following temporary values:

i: Current vertex coordinate index, initially 0.  
cp: Current position, initially (0, 0).  
ct: Current tangent, initially undefined.  
cs: Pending cap style, initially butt.  
sp: Start position, initially undefined.  
st: Start tangent, initially undefined.

Each command is processed as follows, depending on its type:

**MOVE\_TO\_NV:**

- \* If <ct> is defined, draw a butt end cap at <cp>.
- \* cp = (c[i+0], c[i+1]), ct = undefined, cs = butt, i += 2.

**LINE\_TO\_NV, QUADRATIC\_BEZIER\_TO\_NV, and CUBIC\_BEZIER\_TO\_NV:**

- \* Draw the segment corresponding to the command type.
- \* If <ct> is undefined, draw a start cap of style <cs> at <cp>.
- \* If <ct> is defined, draw a join at <cp> between <ct> and the start tangent of the segment.
- \* If the previous command is START\_MARKER\_NV, replace <st> with the start tangent of the segment.
- \* cp = end point, ct = end tangent, i += num.

**START\_MARKER\_NV:**

- \* If <ct> is defined, draw a butt end cap at <cp>.
- \* ct = undefined, cs = butt, sp = cp, st = undefined.

**CLOSE\_NV:**

- \* If <sp> is undefined, ignore the command.
- \* Draw a straight line from <cp> to <sp>.
- \* If <ct> is undefined, draw a start cap of style <cs> at <cp>.
- \* If <ct> is defined, draw a join at <cp> between <ct> and the direction of the line.
- \* If <st> is undefined, draw a butt end cap at <sp>.
- \* If <st> is defined, draw a join at sp between the direction of the line and <st>.
- \* cp = sp, ct = undefined, cs = butt.

**STROKE\_CAP[0-3]\_NV:**

- \* If <ct> is defined, draw an end cap of the specified style.
- \* ct = undefined, cs = style specified by the command.

**Path programs**

Paths are drawn using a special type of program object called a path program. Path programs function like normal program objects, except that

they do not allow a vertex shader to be specified. Path programs are created with a new function `CreatePathProgramNV()`.

#### Fragment shader

Fragment shader depth values are obtained by transforming the homogenous vertex coordinates (x, y, 0, 1) into the clip space. This enables mixing 3D and path geometry using depth buffering.

Since path programs do not support vertex shaders, path fragment shaders cannot make use of user-defined varyings. Instead, this extension adds built-in variables `gl_PathCoord[0-3]` of type `vec4` that receive interpolated vertex positions transformed with their respective `MATRIX_PATH_COORD[0-3]_NV` matrices.

The value of `gl_FrontFacing` is undefined when rendering paths.

#### The rest of the pipeline

When rendering paths, stencil functionality and backface culling are not applied. Blending, dithering, depth test, scissor test, polygon offset, and multisampling are applied as with other primitives.

Even though stencil test and operation are unavailable when rendering paths, the original contents of the stencil buffer are retained.

#### Path buffers

Path buffers facilitate efficient rendering of animated text or other instanced path geometry by making it possible to render multiple path objects with a single draw call. A path buffer contains a list of path object handles and associated translation vectors.

#### Invariance rules

Changing path parameters, viewport, transformations and clipping parameters may result in a different set of pixels to be rendered.

#### New Procedures and Functions

```
uint CreatePathNV(          enum datatype,
                           sizei numCommands,
                           const ubyte* commands );

void DeletePathNV(          uint path );

void PathVerticesNV(        uint path,
                           const void* vertices );

void PathParameterfNV(      uint path,
                           enum paramType,
                           float param );

void PathParameteriNV(      uint path,
                           enum paramType,
                           int param );

uint CreatePathProgramNV(    void );

void PathMatrixNV(          enum target,
```

```

const float* value );

void DrawPathNV(      uint path,
                      enum mode );

uint CreatePathbufferNV(  sizei capacity );

void DeletePathbufferNV(  uint buffer );

void PathbufferPathNV(   uint buffer,
                        int index,
                        uint path );

void PathbufferPositionNV( uint buffer,
                          int index,
                          float x,
                          float y );

void DrawPathbufferNV(   uint buffer,
                        enum mode );

```

#### New Types

None

#### New Tokens

Accepted as the <paramType> parameter of PathParameterNV:

PATH_QUALITY_NV	0x8ED8
FILL_RULE_NV	0x8ED9
STROKE_CAP0_STYLE_NV	0x8EE0
STROKE_CAP1_STYLE_NV	0x8EE1
STROKE_CAP2_STYLE_NV	0x8EE2
STROKE_CAP3_STYLE_NV	0x8EE3
STROKE_JOIN_STYLE_NV	0x8EE8
STROKE_MITER_LIMIT_NV	0x8EE9

Values for the ILL\_RULE\_NV path parameter:

EVEN_ODD_NV	0x8EF0
NON_ZERO_NV	0x8EF1

Values for the CAP[0-3]\_STYLE\_NV path parameter:

CAP_BUTT_NV	0x8EF4
CAP_ROUND_NV	0x8EF5
CAP_SQUARE_NV	0x8EF6
CAP_TRIANGLE_NV	0x8EF7

Values for the JOIN\_STYLE\_NV path parameter:

JOIN_MITER_NV	0x8EFC
JOIN_ROUND_NV	0x8EFD
JOIN_BEVEL_NV	0x8EFE
JOIN_CLIPPED_MITER_NV	0x8EFF

Accepted as the <target> parameter of PathMatrixNV:

MATRIX_PATH_TO_CLIP_NV	0x8F04
------------------------	--------

MATRIX_STROKE_TO_PATH_NV	0x8F05
MATRIX_PATH_COORD0_NV	0x8F08
MATRIX_PATH_COORD1_NV	0x8F09
MATRIX_PATH_COORD2_NV	0x8F0A
MATRIX_PATH_COORD3_NV	0x8F0B

Accepted as the <mode> parameter of DrawPathbufferNV:

FILL_PATH_NV	0x8F18
STROKE_PATH_NV	0x8F19

Accepted as path commands by CreatePathNV:

MOVE_TO_NV	0x00
LINE_TO_NV	0x01
QUADRATIC_BEZIER_TO_NV	0x02
CUBIC_BEZIER_TO_NV	0x03
START_MARKER_NV	0x20
CLOSE_NV	0x21
STROKE_CAP0_NV	0x40
STROKE_CAP1_NV	0x41
STROKE_CAP2_NV	0x42
STROKE_CAP3_NV	0x43

#### Additions to Chapter 2 of the OpenGL ES Specification

Add the following error conditions to Chapter 2.8, under DrawArrays:

"An INVALID\_OPERATION error is generated if the current program is a path program."

Add the following error conditions to Chapter 2.8, under DrawElements:

"An INVALID\_OPERATION error is generated if the current program is a path program."

Add the following error conditions to Chapter 2.15, under AttachShader.

"An INVALID\_OPERATION error is generated if the program is a path program and the shader is a vertex shader."

Add the following error conditions to Chapter 2.15, under LinkProgram.

"Linking a program without a vertex shader will not fail if the program is a path program."

#### Additions to Chapter 3 of the OpenGL ES Specification

Add a new section between sections 3.5 (Polygons) and 3.6 (Pixel Rectangles)

##### "3.6 Paths

This extension adds a new type of primitive, paths, to OpenGL ES' primitives - points, lines, polygons, pixel rectangles and bitmaps.

##### 3.6.1 Path objects

New path objects are created with the call



```

uint CreatePathNV( enum datatype,
                   sizei numCommands,
                   const ubyte* commands );

```

where <datatype> is the vertex data type and it must be one of [UNSIGNED\_]BYTE, [UNSIGNED\_]SHORT, [UNSIGNED\_]INT, FLOAT, FIXED, <numCommands> is the number of commands in the path definition and <commands> is a pointer to an unsigned byte array of commands. Valid commands are listed below. The function returns a non-zero handle to the object or 0 on error.

An INVALID\_ENUM error is generated if <datatype> is not one of the values specified above. An INVALID\_VALUE error is generated if <numCommands> is less than zero or <numCommands> is greater than zero and <commands> is NULL or the <commands> array contains an invalid command.

TODO note that path objects can be shared between multiple contexts

Path objects are deleted with the command

```

void DeletePathNV( uint path );

```

where <path> is the handle to the path object to delete. If the path is assigned to one or more path buffers, path resources are freed only when the last reference to the path is removed. Path handle is invalid after a call to DeletePathNV. An INVALID\_VALUE error is generated if the path object does not exist.

Path vertices are specified with the command

```

void PathVerticesNV( uint path,
                    const void* vertices );

```

where <path> is the handle to the path object and <vertices> is a pointer to an array of vertices. <vertices> must contain at least as many coordinate tuples as is consumed by the associated path commands, otherwise the results are undefined, and may lead to a program crash. If <vertices> contains more coordinates than consumed by the path commands, the rest are silently ignored.

An INVALID\_VALUE error is generated if the specified <path> object does not exist or if <vertices> is NULL and the path command requires vertices.

Path parameters are set using the commands

```

void PathParameterfNV( uint path,
                      enum paramType,
                      float param );

void PathParameteriNV( uint path,
                      enum paramType,
                      int param );

```

where <path> is the path object handle, <paramType> is the parameter to set and <param> is the value of the parameter.

The following symbols are accepted as <paramType>:

PATH\_QUALITY\_NV

Maximum allowed deviation from the ideal path measured in pixels. The

default value is 0.5 pixels.

**FILL\_RULE\_NV**  
 Fill rule to use for filling paths. <param> must be either EVEN\_ODD\_NV or NON\_ZERO\_NV. The default value is EVEN\_ODD\_NV.

**STROKE\_CAPn\_STYLE\_NV**  
 cap style for the cap index n used when stroking a path. The default values are CAP\_BUTT\_NV.

**STROKE\_JOIN\_STYLE\_NV**  
 join style used when stroking a path. The default value is JOIN\_MITER\_NV.

**STROKE\_MITER\_LIMIT\_NV**  
 miter limit used when stroking a path with miter joins. If a join angle exceeds the limit, a miter join is converted into a bevel join. The default value is 4.

If paramType is PATH\_QUALITY\_NV in PathParameteriNV(), param is converted to a float. If paramType is not PATH\_QUALITY\_NV in PathParameterfNV(), param is converted to an int.

An INVALID\_VALUE error is generated if the <path> object does not exist. An INVALID\_ENUM error is generated if <paramType> is not any of the above. An INVALID\_VALUE error is generated if <paramType> is PATH\_QUALITY\_NV and <param> <= 0 or <paramType> is STROKE\_MITER\_LIMIT\_NV and <param> < 1. An INVALID\_ENUM error is generated if <paramType> is FILL\_RULE\_NV and param is not a valid fill rule, <paramType> is STROKE\_CAPn\_STYLE\_NV and <param> is not a valid cap style, or <paramType> is STROKE\_JOIN\_STYLE\_NV and <param> is not a valid join style.

Path transformations are set using the call

```
void PathMatrixNV( enum target,
                  const float* value );
```

where <value> must specify a 4x4 matrix. The following values are accepted as the <target> parameter:

**MATRIX\_PATH\_TO\_CLIP\_NV**  
 used for transforming path vertices into clip space when drawing a path.

**MATRIX\_STROKE\_TO\_PATH\_NV**  
 used for transforming the pen when stroking a path. The vertices are subsequently transformed into clip space by MATRIX\_PATH\_TO\_CLIP\_NV matrix. Only the top-left 2x2 submatrix is used.

**MATRIX\_PATH\_COORDn\_NV**  
 used for generating values for gl\_PathCoord[0-3] varyings for fragment shader by transforming vertex positions.

The default value for all matrices is the identity matrix.

An INVALID\_ENUM error is generated if <target> is not any of the above. An INVALID\_VALUE error is generated if value is NULL.

A path is rendered using the call

```
void DrawPathNV( uint path,
                enum mode );
```

where <path> is the path to be drawn and <mode> must be either FILL\_PATH\_NV or STROKE\_PATH\_NV.

An `INVALID_VALUE` error is generated if the `<path>` does not exist. An `INVALID_OPERATION` error is generated if there is no current program, the current program is not a path program, stencil test is enabled, polygon mode is not `GL_FILL`, or shade model is not `GL_SMOOTH`. An `INVALID_ENUM` is generated if `<mode>` is not `FILL_PATH_NV` or `STROKE_PATH_NV`.

### 3.6.2 Path programs

A path program is a special type of program object that otherwise behaves like a normal program object, but allows attaching only a fragment shader. Path programs are created using the command

```
uint CreatePathProgramNV( void );
```

The function returns 0 on error (i.e. `OUT_OF_MEMORY`).

TODO describe `LinkProgram` error conditions here for clarity?

### 3.6.3 Path buffers

Path buffers can be used for efficiently rendering multiple instances of a set of path objects with a single draw call. Each path in a path buffer has an associated position vector that allows specifying a model space position offset for that path. Path buffers are created using the function

```
uint CreatePathbufferNV( sizei capacity );
```

where `<capacity>` is the number of paths in a path buffer. This function returns a non-zero handle to a path buffer object or 0 on error.

An `INVALID_VALUE` error is generated if `capacity < 0`.

TODO note that path buffer objects can be shared between multiple contexts

Path buffers are deleted using the call

```
void DeletePathbufferNV( uint buffer );
```

where `<buffer>` is the handle to the path buffer. Path buffer handle is invalid after a call to `DeletePathbufferNV`.

An `INVALID_VALUE` error is generated if the path buffer does not exist.

A path can be added to or removed from a path buffer with the function

```
void PathbufferPathNV( uint buffer,  
                      int index,  
                      uint path );
```

where `<buffer>` is the path buffer object handle, `<index>` is the index of the path buffer slot and `<path>` is the path object handle. Path buffer paths are mutable and can be re-specified later. Calling `PathbufferPathNV` with `<path>` set to zero removes path from the path buffer and leaves the slot corresponding to `<index>` empty.

An `INVALID_VALUE` error is generated if the path buffer object `<buffer>` does not exist, or `<index>` is less than zero, or greater than or equal to the path buffer capacity.

Path buffer path position vector is specified using the call

```
void PathbufferPositionNV( uint buffer,
                           int index,
                           float x,
                           float y );
```

where <buffer> is the path buffer object handle, <index> is the index of the path buffer slot and <x> and <y> specify the translation. Path buffer path translations are mutable and can be re-specified later.

An INVALID\_VALUE error is generated if the path buffer does not exist or index < 0 or index >= path buffer capacity.

All paths in a path buffer are rendered using the command

```
void DrawPathbufferNV( uint buffer,
                       enum mode );
```

An INVALID\_VALUE error is generated if <buffer> does not exist. An INVALID\_OPERATION error is generated if there is no current program, the current program is not a path program, stencil test is enabled, polygon mode is not GL\_FILL, or shade model is not GL\_SMOOTH. An INVALID\_ENUM is generated if <mode> is not FILL\_PATH\_NV or STROKE\_PATH\_NV. The effect of a DrawPathbufferNV call is the same as if DrawPathNV was called for each individual path reference in the path buffer, ordered from the first index to the last."

Add the following to Chapter 3.11 in the section Shader Inputs.

"The value of gl\_FrontFacing is undefined if the current program is a path program."

"If the current program is a path program and fragment shader has defined varying variables gl\_PathCoord[0-3], they will receive interpolated vertex coordinates transformed with their respective MATRIX\_PATH\_COORD[0-3]\_NV".

Additions to Chapter 4 of the OpenGL ES Specification

Add the following to the end of Chapter 4.1.5 Stencil Test.

"Stencil functionality is not applied when rendering paths. Path rendering will generate an error if stencil testing is enabled."

Issues

1. Should we use vertex shader or fixed function transform?

RESOLUTION: Introduce minimal fixed function transform and texgen functionality.

DISCUSSION: The problem with vertex shader is that it allows changing vertex/control point positions arbitrarily, but path geometry only makes sense if it remains planar. A similar problem occurs with vertex attributes: since attributes of three vertices define interpolation on a plane, the attributes of the rest of the vertices cannot be chosen freely for the interpolation to remain well-defined. In effect, this forces vertex attributes to be derived from vertex positions, which can be described by a matrix multiplication. Furthermore, many implementations are expected to cache the results of flattening and/or triangulation, which is much simpler in case of fixed function

transform.

Downsides to using fixed function transformation include the lack of support for morphing. We expect most content to not use morphing, so a viable alternative is modifying the path coordinates. Another downside is that we're reintroducing fixed function transform stage into ES2. In our opinion, the downsides of adopting vertex shader solution outweigh this concern.

2. Should we leverage VBOs/vertex arrays for path coordinates and commands?

RESOLUTION: Involving VBOs would make the extension much messier.

3. Should we have elliptical arc segments?

RESOLUTION: Not in this version. It is fairly straightforward to convert arcs into quadratic beziers in an application when content is loaded.

4. Should we have vertex indices?

RESOLUTION: No. This would unnecessarily complicate the API.

5. Should we use 1 - 4 component vectors as vertex position?

RESOLUTION: No. It is not clear how non-planar geometry would be rendered.

6. Should we support perspective transformations?

RESOLUTION: Yes.

7. Should we support stroking?

RESOLUTION: Yes, all vector graphics formats have stroking. Stroking is a rather involved process both implementation-wise and computationally, so it is a good candidate for being implemented in a driver.

8. Should we allow modification of paths?

RESOLUTION: Modifying coordinates is needed for animation, especially since we ignore vertex shader. There is no good use case for modifying commands.

9. What happens in case a contour is not closed?

RESOLUTION: The result is undefined. Alternatively we could have an error check, but that is an extra burden for implementations and extra work in the common case where the path data is ok.

10. Should we support edge antialiasing?

RESOLUTION: No. ES2 doesn't support edge antialiasing for other primitives either, and this extension is compatible with multisampling.

11. What invariance requirements should we impose?

SUGGESTION:

- 1) Rendering the same path with the same state must generate the same

pixels. This is the normal GL invariance requirement.

- 2) If two adjoining paths have a shared curve defined by exactly the same vertices (bitwise exact), there can be no gaps. The curve direction can change and the invariant must still hold.
- 3) UNRESOLVED: If two paths have a shared curve, plus one of the paths has extra geometry that intersects the shared curve, does the no-gap requirement still have to hold? This is problematic for implementations that generate extra vertices at intersection points in the process of triangulation/path simplification.

#### Revision History

#0.11 - 2008/09/16 - Tero Karras  
#0.10 - 2008/09/15 - Tero Karras  
#0.9 - 2008/09/12 - Jussi Rasanen

## NV\_system\_time

### Name

NV\_system\_time

### Name Strings

EGL\_NV\_system\_time

### Contact

Jason Allen, NVIDIA Corporation (jallen 'at' nvidia.com)

### Dependencies

Requires EGL 1.2

### Overview

This extension exposes an alternative method of querying the system time from the driver instead of the operating system.

### Issues

Add 64 bit types?

Yes, EGL doesn't support any 64 bit types so this extension adds int64 and uint64 types.

### New Types

```
typedef long int EGLint64NV;  
typedef unsigned long int EGLuint64NV;
```

### New Procedures and Functions

```
uint64NV GetSystemTimeFrequencyNV(void);  
uint64NV GetSystemTimeNV(void);
```

### New Tokens

None

### Description

The command:

```
uint64NV GetSystemTimeFrequencyNV(void);
```

returns the frequency of the system timer, in counts per second. The frequency will not change while the system is running.

The command:

```
uint64NV GetSystemTimeNV(void);
```

returns the current value of the system timer. The system time in seconds can be calculated by dividing the returned value by the frequency returned by the GetSystemTimeFrequencyNV command.

Multiple calls to `GetSystemTimeNV` may return the same values, applications need to be careful to avoid divide by zero errors when using the interval calculated from successive `GetSystemTimeNV` calls.

#### Usage Example

```
EGLuint64NV frequency = eglGetSystemTimeFrequencyNV();

loop
{
    EGLuint64NV start = eglGetSystemTimeNV() / frequency;

    // draw

    EGLuint64NV end = eglGetSystemTimeNV() / frequency;

    EGLuint64NV interval = end - start;
    if (interval > 0)
        update_animation(interval);

    eglSwapBuffers(dpy, surface);
}
```



**Notice**

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication or otherwise under any patent or patent rights of NVIDIA Corporation.

Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all information previously supplied. NVIDIA Corporation products are not authorized for use as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

**Trademarks**

NVIDIA, the NVIDIA logo, Tegra, GeForce, NVIDIA Quadro, and NVIDIA CUDA are trademarks or registered trademarks of NVIDIA Corporation in the United States and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

**Copyright**

© 2008-2010 NVIDIA Corporation. All rights reserved.

**NVIDIA.**

NVIDIA Corporation

2701 San Tomas Expressway

Santa Clara, CA 95050

[www.nvidia.com](http://www.nvidia.com)